

MTDF Validation Suite Appendix: Reproducibility Index and Test Outcomes

Ingo Mesche
Independent Researcher, Malta
imesche@outlook.com

V75 / July 2026

Shared appendix referenced by MTDF Companion Papers Parts I to IV. Updated for Validation_Dashboard_V75.html.

Abstract

This appendix consolidates the reproducibility index, test scripts, and summary diagnostics for the MTDF validation suite. It provides the complete inventory of tests referenced across the companion papers (MTDF_01 through MTDF_08), including script locations, input data, and pass/fail outcomes under the V18 strict protocol.

Claim status. This paper distinguishes between: (i) calibration anchors, (ii) validation targets, (iii) diagnostic consistency checks, and (iv) speculative extensions. Only results explicitly labelled as validation targets are used as evidence for the core MTDF claim. Diagnostic and speculative sections are included to define future tests, not to claim established confirmation.

Conventions. Growth and S8 material defers to the sealed v2 chains (the v1 Phase-5 report is retired as a gauge artifact; see MTDF_01 Section 7.3 and MTDF_07 Section 6).

1 Purpose and how to use this appendix

This appendix consolidates the reproducibility index and quick-look outcomes for the test scripts used across the MTDF companion papers. The intent is simple: keep the core papers readable, while keeping the test inventory, scripts, and summary diagnostics in one stable place.

Files referenced here:

- Companion Part I: Environmental and local evidence
- Companion Part II: Gravity sector and lensing validation (*MTDF_03* [see companion paper])
- Companion Part III: Photon coupling, redshift, and early universe links
- Companion Part IV: Cosmological validation and high-energy test programme
- Interactive dashboard: [Validation_Dashboard_V75.html](#)

Scope note: This appendix is a reproducibility index. It does not change the evidential status of any result. Claim status is defined in the corresponding main or companion paper. This appendix records scripts and diagnostics. Statistical weight remains with the distance-calibrated observables and the calibration versus validation separation defined in the dashboard workflow.

2 Test index and reproducibility

This appendix records the scripts used to generate the figures and summary diagnostics for Parts I to IV of the broader MTDF validation effort. The SN Ia $\times$ void result is primary evidence within Companion Part I’s extension programme; it is not part of the V18 strict validation core. The remaining tests are included as constraints, consistency checks, and forward-looking discriminants across Parts II and III.

2.1 Script index by paper component

Paper component	Test	Script	Role
Part I (environment and independent distances)	SN Ia $\times$ void analysis	sn_void_*.py	Primary evidence for an environment-dependent distance-calibrated residual.
	Zero-velocity shell	mtdf_zero_velocity_shell_test.py	Null and pipeline integrity check.
	Lensed image Δz	mtdf_lensing_redshift_split_test.py	Path-dependence test using multiple lensed images of the same source.
Part III (photon coupling κ) <i>[speculative]</i>	Achromaticity	mtdf_achromaticity_test.py	Wavelength independence requirement for any stress-induced redshift.
	Dipole mismatch	mtdf_dipole_mismatch_test.py	Checks whether a stress-induced dipole can contribute to the quasar dipole anomaly.
	Dipole alignment	mtdf_dipole_vector_alignment.py	Directional consistency check between excess dipole and density gradient.
	Distance duality	mtdf_distance_duality_test.py	Etherington relation and potential environment split.
	Sandage-Loeb drift	mtdf_sandage_loeb_drift_test.py	Long-baseline, future definitive test of evolving redshift physics.
	Synthesis plot	mtdf_synthesis_plot.py	Conceptual summary figure linking multiple anomalies to the same scaling law.
	S8 tension and growth	mtdf_growth_S8_analysis.py	Large-scale structure growth consistency.
Part IV (cosmological and lensing)	BAO void split	mtdf_bao_void_split_test.py	Environment split for BAO-scale observables.
	GW standard sirens	mtdf_gw_siren_test.py	Independent distance ladder cross-check using gravitational waves.
	Lensing versus RSD (E_G)	mtdf_lensing_rsd_test.py	Modified-gravity consistency test, not a direct probe of photon coupling.
	Cluster mass MDR	mtdf_cluster_mass_discrepancy_test.py	Compares dynamical and lensing mass proxies in clusters.
	AP void shape	mtdf_ap_void_test.py	Alcock-Paczynski void geometry constraint.
Part II (gravity sector and lensing)	Strong lensing time delays	step20_strong_lensing.py	Time-delay distance consistency check under $\eta = 1$.
	Elliptical galaxy dispersions	step22_elliptical_dispersions.py	Pressure-supported system check (SLACS-style).
	Satellite kinematics	step22b_satellite_kinematics.py	Digitised and table-based comparison (More+2011 style) with robustness notes.

2.2 Outcome snapshot

Script	Headline outcome	Interpretation at current sensitivity
mtdf_zero_velocity_shell_test.py	Consistent	No significant artifact detected in a designed null scenario.
mtdf_distance_duality_test.py	Consistent	No significant violation of $\eta = 1$ at current precision; environment split remains a future target.
mtdf_lensing_redshift_split_test.py	Consistent	No significant Δz between multiple lensed images; current spectroscopy is above the $\kappa \lesssim 10^{-3}$ signal level.
mtdf_achromaticity_test.py	Consistent at current precision	No evidence for chromatic redshift; any measured offsets are likely systematics at the present 10^{-5} level.
mtdf_dipole_mismatch_test.py	Suggestive (partial amplitude)	At $\kappa \approx 0.001$, MTDF reproduces the correct sign and roughly 42% of the reported quasar dipole excess amplitude.
mtdf_dipole_vector_alignment.py	Moderate directional agreement	Alignment is not tight enough to be a strong constraint, but it remains informative for future catalogues.
mtdf_sandage_loeb_drift_test.py	Future-only	Predicted environment split in $\dot{z}$ is far below current detection limits for the derived κ value ($f_{\text{kick}}/3$).
mtdf_synthesis_plot.py	Key synthesis figure	Heuristic summary connecting multiple anomalies across environments. It is not, by itself, a measurement of κ .
mtdf_lensing_rsd_test.py	Consistent	E_G data are broadly consistent with GR at current precision; not a direct κ probe.
mtdf_cluster_mass_discrepancy_test.py	Historical exploratory row; under the frozen architecture $\eta = 1$ gives $M_{\text{dyn}}/M_{\text{lens}} = 1.000$ exactly (see MTDF_07, Test C5): this channel is not a discriminator.	Historical exploratory note; superseded by the $\eta = 1$ result (MTDF_07, Test C5).
mtdf_ap_void_test.py	Consistent	Void shape residuals consistent with zero after RSD correction; κ effects are below current precision.
mtdf_growth_S8_analysis.py	See Part IV (Phase 5)	Growth and S8 are treated in Part IV via the sealed v2 full-Planck chains (the v1 Phase-5 report is retired as a gauge artifact); this script is an earlier exploratory check and its results are not used there.
mtdf_bao_void_split_test.py	Exploratory, not in the papers	Exploratory diagnostic; BAO environment splits are not part of the validated core and are not presented in Part IV.
mtdf_gw_siren_test.py	Exploratory, not in the papers	Exploratory diagnostic; GW standard-siren distance checks are not presented in the papers and carry no evidential weight here.
step20_strong_lensing.py	See Part II	Time-delay distance $\chi^2/\nu = 0.77$; isothermal slope $\gamma = 2.000$ (structural prediction). See gravity-sector companion.
step22_elliptical_dispersions.py	See Part II	Mean velocity dispersion bias -1.4%; $\chi^2/\nu = 1.81$ (gas-corrected). See gravity-sector companion.
step22b_satellite_kinematics.py	See Part II	More+2011 amplitude test: mean $ \text{bias} $ 6.3%; profile $\chi^2/\nu = 3.60$ (best nuisance). See gravity-sector companion.

2.3 Synthesis figure

Figure A.1 provides a compact visual summary that places three anomalies on a single stress-scaling trend. This figure is intended as a high-level synthesis and motivates a unified search strategy. It should not be interpreted as a standalone parameter inference for κ .




Figure 1.

Figure A.1. The stress-redshift main sequence synthesis plot (generated by `mtdf_synthesis_plot.py`).

2.4 P1-SPARC test definition

Statistic. The P1-SPARC benchmark reports the aggregate root-mean-square of the log-space residual $\log_{10}(v_{\text{obs}}/v_{\text{model}})$ over the full SPARC sample (Lelli et al. 2016, AJ 152, 157; 175 galaxies). The achieved value is 0.185 dex against a target of 0.174 ± 0.011 dex, below the 0.25 dex reference threshold. This is an aggregate amplitude statistic, not a per-galaxy rotation-curve shape fit, and no per-galaxy χ^2 of zero is claimed.

Model and parameters. The MTDf velocity is $v_{\text{model}}(r) = v_{\text{bar}}(r) \sqrt{1 + \alpha/(1 + r/\beta)}$ with $\alpha = 1.30$ and $\beta = 22.7$ Mpc frozen. Because β greatly exceeds galactic radii, the enhancement is effectively a constant factor $\sqrt{1+\alpha} \approx 1.52$ in velocity (a factor $1+\alpha \approx 2.3$ in acceleration) across the rotation curve.

Mass-to-light treatment and error model. v_{bar} is built from the SPARC mass-model decomposition (gas, disk, bulge); the aggregate scatter is evaluated at the mass-model normalisation, and the residual uses the observed velocity uncertainties. For reference, a per-galaxy shape fit at the fixed SPARC-standard mass-to-light ratios (0.5 for disks, 0.7 for bulges at $3.6 \mu\text{m}$) yields per-galaxy reduced chi-squared of order 10–280 (e.g. NGC 3198 ≈ 72 , NGC 2403 ≈ 57 , NGC 7331 ≈ 281), confirming that the constant enhancement reproduces the amplitude of the rotation curves but not their shape.

Shape diagnostic (radial acceleration relation). On the radial acceleration relation, the observed data follow a low-acceleration log-log slope of ≈ 0.58 (consistent with the McGaugh et al. 2016 value near 0.5), whereas a constant enhancement predicts a slope of 1.0 (a parallel offset). MTDf therefore matches the approximate amplitude of the relation (median $g_{\text{obs}}/g_{\text{bar}} \approx 2.7$ against a predicted $1 + \alpha = 2.3$) but not its curvature; the constant-enhancement parameterisation used here does not bend. The curvature is produced by the low-Q branch of the frozen unified action (MTDF_01, Section 2A), which resolves the shape-level question; this appendix’s parameterisation is retained as the historical effective description.

Scope. The galaxy-scale pillars are treated as scalar benchmarks and are excluded from the V18 strict scoring totals, so the shape-level limitation above does not affect the strict vector results.

Audit artefacts (per-galaxy table and the radial-acceleration-relation diagnostic) are provided in the reproducibility package under `validation/output/audit_sparc/`.

3 Data and code availability

The complete MTDF reproducibility package, including the modified CLASS solver (`class_mtdf`), the V19 strict validation workbook, the `run_validate.py` pipeline, the gravity-sector artefact manifest, and the Phase 2 CosmoPower emulator cache, is archived on Zenodo (concept DOI 10.5281/zenodo.19741058; this release, v1.1.7, DOI 10.5281/zenodo.21525855) and mirrored on GitHub at github.com/IMesche/MTDF. The v1.1.7 Git tag corresponds to the archived Zenodo release. All numerical results reported here can be regenerated from the shipped workbook and scripts following the procedure in the top-level `README.md`.