

MTDF Independent GPU Validation: Phases 1–6

Ingo Mesche

Independent Researcher, Malta

imesche@outlook.com

V74 / February 2026

Blind reproduction, CMB consistency, environment signal confirmation, sensitivity forecasts, full Planck hard-falsifier test, and cross-probe discriminators

Hardware: NVIDIA RTX 4080 SUPER (16 GB VRAM, CUDA 12.9)

Abstract. This document reports the results of a fully independent, GPU-accelerated validation of Mesche’s Tensor Dynamics Framework (MTDF). Six phases were executed over multiple sessions using code written from scratch, with no shared libraries or routines from the original MTDF analysis pipeline. Phase 1 reproduces all V18 dashboard χ^2 values to four decimal places. Phase 2 tests MTDF against the Planck 2018 CMB power spectrum via MCMC with a perturbative correction layer, finding $\Delta\chi^2 < 1$ and an unconstrained early-field-energy amplitude. Phase 3 independently confirms the $\text{SN} \times \text{void}$ environment signal reported in MTDF_02, with bootstrap confidence intervals excluding zero. Phase 4 forecasts that 5σ detection of the environment signal requires approximately 3,400 low-redshift supernovae, placing definitive confirmation within the reach of LSST/Rubin. Phase 5 executes the Priority 1 hard falsifier: the full Planck 2018 plik TTTEEE + lowl + lensing likelihood using the `class_mtdf` Boltzmann solver with modified perturbation equations, finding $\Delta\chi^2 = +0.63$. This MTDF implementation is statistically indistinguishable from ΛCDM at the CMB within the precision and assumptions of this likelihood comparison. Phase 6 extends validation to cross-probe discriminators: a redshift transition scan (3.6σ at $z < 0.03$), weak lensing $\times$ environment (KiDS-1000; systematics-clean null with explicit upper limits), derived parameter consistency (S_8 tension reduced), and CMB lensing $\times$ voids (Planck PR4 stacking on DESI BGS and BOSS DR12 void catalogues; pipeline validated against published detections).

Contents

1	Scope and methodology	3
1.1	Hardware and software	4
2	Phase 1: Independent χ^2 reproduction	4
2.1	Strict total	4
2.2	Vector pillar breakdown	4
3	Phase 2: Planck CMB consistency	4
3.1	Emulator validation	5
3.2	Fixed-parameter comparison	5
3.3	Normalisation correction	5
3.4	Production MCMC results	5
3.4.1	Model comparison	5
3.4.2	k_f posterior	6

3.4.3	H_0 posterior	6
3.4.4	Per-probe breakdown (combined narrow best-fit)	6
4	Phase 3: SN $\times$ void environment signal	6
4.1	Environment coefficient γ_{env}	6
4.2	NGC/SGC split	7
4.3	Survey fixed effects	7
4.4	Leave-one-survey-out (LOSO)	7
4.5	Redshift modulation: the signature result	7
4.6	Non-parametric confirmation	7
4.6.1	Permutation test (10,000 shuffles)	7
4.6.2	Block bootstrap (5,000 resamples), 95% CI	7
4.7	Numerical integrity	7
5	Phase 4: Sensitivity forecasts	8
5.1	Key structural insight	8
5.2	Detection threshold table	8
5.3	Interpretation	8
6	5b. Phase 3b: Asymmetry and detectability diagnostics	9
6.1	5b.1 Test A: IPW matched redshift	9
6.2	5b.2 Test B: Void-SN joint density mapping	9
6.3	5b.3 Test C: Wald test and parametric bootstrap	9
6.4	5b.4 Test D: Geometry and mode-coupling check	9
6.5	5b.5 Test E: Null signal injection (SGC recovery)	9
6.6	5b.6 Interpretation	10
7	Phase 5: Full Planck plik + class_mtdf	10
7.1	χ^2 breakdown by likelihood component	11
7.2	Model selection	11
7.3	Best-fit cosmological parameters	11
7.4	Interpretation	11
7.5	Planck degeneracy structure and posterior stability (MCMC converged)	12
7.5.1	Convergence statistics	12
7.5.2	Posterior parameter comparison (68% CI)	12
7.5.3	k_f posterior	12
7.5.4	k_f identifiability diagnostics	12
7.5.5	Best-fit robustness across Planck subsets	13
7.5.6	Phase 5 robustness checks (Tests 2–4)	13
8	6b. Phase 6: Cross-probe discriminator tests	14
8.1	Summary table	14
8.2	6b.1 Test A: Redshift transition scan	15
8.3	6b.2 Test B: Weak lensing $\times$ environment (void centres)	15
8.4	6b.3 Test B2: Trough and ridge lensing (projected density split)	15
8.5	6b.4 Test C: Derived parameter consistency	15
8.6	6b.5 Test D: CMB lensing $\times$ DESI BGS voids (low redshift)	16
8.7	6b.6 Test E: CMB lensing $\times$ BOSS DR12 voids (benchmark)	16
8.8	6b.7 Test C2: S_8 cross-probe coherence	16
8.9	6b.8 Test C3: BAO residual structure	17
8.10	6b.9 Test C4: $f\sigma_8 \times$ environment	17
8.11	6b.10 Test C5/C5b: Cluster $M_{\text{dyn}}/M_{\text{lens}}$ - Resolved	17
9	Synthesis and framework narrative	18

9.1	Summary of results (17/17 complete)	18
9.2	Claim status taxonomy	19
9.2.1	Status by claim	20
9.3	Framework narrative	21
9.4	What survives intact	21
9.5	Joint implications for late-time cosmological tensions	22
9.6	Paper implications	22
10	Open questions and next steps	22
10.1	NGC/SGC asymmetry diagnostics (Phase 3b, complete)	22
10.2	Falsifier and discriminator tests	23
11	Appendix: Technical notes	23
11.1	A.1 k_f normalisation conventions	23
11.2	A.2 V18 parameter reference	24
11.3	A.3 Historical context: November 2025 runs	24
12	References	24

Claim status. This paper distinguishes between: (i) calibration anchors, (ii) validation targets, (iii) diagnostic consistency checks, and (iv) speculative extensions. Only results explicitly labelled as validation targets are used as evidence for the core MTDf claim. Diagnostic and speculative sections are included to define future tests, not to claim established confirmation.

1 Scope and methodology

The purpose of this validation exercise is to subject MTDf to a battery of independent tests that are entirely separate from the original analysis pipeline. All code was written from scratch across multiple sessions. No libraries, routines, or intermediate results from the original MTDf codebase were used for Phases 1, 3, 4, or 6. Phase 2 used the CosmoPower Λ CDM emulator with a bespoke MTDf correction layer. Phase 5 used the existing `class_mtdf` modified Boltzmann solver through cobaya with the full Planck likelihood stack.

The six phases test qualitatively different aspects of the framework:

- **Phase 1** tests internal consistency: can the published χ^2 values be reproduced from raw data using independent code?
- **Phase 2** tests CMB compatibility (perturbative): is MTDf allowed by a compressed CMB likelihood?
- **Phase 3** tests the primary empirical prediction: is the $\text{SN} \times \text{void}$ environment signal real and reproducible?
- **Phase 4** tests future viability: what survey specifications are needed for definitive detection?
- **Phase 5** tests the hard falsifier: does MTDf survive the full Planck likelihood with modified Boltzmann equations?
- **Phase 6** tests cross-probe discriminators: does the transition signature survive in weak lensing, CMB lensing, and derived parameters?

Independence guarantee. Phases 1, 3, and 6 use fully independent implementations. Phase 2 uses a perturbative correction layer on top of a public Λ CDM emulator (CosmoPower), parameterised by a continuous coupling k_f where $k_f = 0$ is Λ CDM and $k_f = 1$ is full MTDf. Phase 4 builds on Phase 3 infrastructure with additional Monte Carlo machinery. Phase 5 uses the

existing `class_mtdf` Boltzmann solver with `cobaya` and the full Planck `plik` likelihood. This is the only phase that uses pre-existing MTDF code, which is appropriate since the test is specifically to validate that code against Planck. Phase 6 uses public data products (KiDS-1000 shear catalogues, Planck PR4 lensing maps, DESI/BOSS void catalogues) with all stacking and analysis code written from scratch. All GPU computations were cross-checked against CPU results.

1.1 Hardware and software

All computations were performed on a desktop workstation equipped with an NVIDIA RTX 4080 SUPER GPU (16 GB VRAM, CUDA 12.9, compute capability 8.9). The software environment consists of Python 3.12 with `cobaya` 3.6, `emcee`, `getdist`, `CuPy`, `JAX`, `TensorFlow`, and `CosmoPower`. All datasets were pre-downloaded: Pantheon+ with full covariance matrix, DESI Y1 BAO, cosmic chronometers, DR16 $f\sigma_8$ growth measurements, and DESI void catalogues (VoidFinder, REVOLVER, VIDE).

2 Phase 1: Independent χ^2 reproduction

PASSED: All V18 dashboard values reproduced to $\delta = 0.000049\text{A}$ fully independent implementation reproduced all V18 dashboard χ^2 values. The code reads raw data files, constructs likelihood functions from scratch, and evaluates MTDF predictions at the V18 parameter point. No dashboard code was consulted or reused.

2.1 Strict total

Metric	Target (V18)	Reproduced	δ
χ^2/ν	1.1683	1.1683	0.000049
Degrees of freedom ν	1745	1745	–

2.2 Vector pillar breakdown

Pillar	χ^2	DOF	χ^2/ν
Pantheon+ SNe Ia	2012.72	1700	1.184
DESI Y1 BAO	15.39	12	1.282
Cosmic chronometers $H(z)$	6.98	15	0.466
DR16 $f\sigma_8$ growth	2.00	3	0.667
CMB distance prior (diagnostic only)	595.33	3	198.44

Plus 15 scalar pillars contributing $\chi^2 = 1.67$, matching exactly.

Significance. Every pillar is reproduced independently, not merely compensating totals. No hidden assumptions, no implementation quirks, no circular calibration detected. The V18 strict protocol is validated by independent audit.

3 Phase 2: Planck CMB consistency

PASSED: $\Delta\chi^2 < 1$; $k_f = 1$ within 95% CI in all runsPhase 2 tests whether MTDF is compatible with the Planck 2018 CMB power spectrum. The approach uses the `CosmoPower` neural network emulator for Λ CDM C_ℓ spectra, augmented by an analytic MTDF correction layer that modifies the sound horizon, angular diameter distance, and ISW contribution as functions of a continuous coupling parameter k_f . The likelihood is the Planck `plik-lite` TTTEEE compressed likelihood (613 bins).

3.1 Emulator validation

Spectrum	Prediction time	Mean error vs CAMB	Max error
TT	7.5 ms	0.09%	0.20%
TE	4.1 ms	0.46%	15.4%*
EE	7.6 ms	0.11%	0.29%

*TE maximum error occurs at zero-crossings where the spectrum passes through zero. This is expected and numerically harmless.

3.2 Fixed-parameter comparison

Cosmology	χ^2	Bins	χ^2/bin
Λ CDM (CAMB reference)	584.45	613	0.953
MTDF at $k_f = 1$ (corrected)	585.26	613	0.955
$\Delta\chi^2$	+0.81		

Λ CDM $\chi^2 = 584.45$ matches the Planck 2018 published value (~ 583). Pipeline correctly wired.

At the full MTDF predicted amplitude ($k_f = 1$), the corrections to the CMB are perturbatively small: the sound horizon shift is -0.071% , the angular scale shift is $+0.021\%$, and the peak shifts are less than one multipole at the first three acoustic peaks. The $\Delta\chi^2$ of $+0.81$ is far below any statistical significance threshold.

3.3 Normalisation correction

During development, a normalisation inconsistency was identified and corrected. Three different k_f conventions existed in the codebase (see Appendix). After correction, the physical mapping is:

$$f_{\text{kick}} = \frac{\lambda_{\text{MTDF}}}{24} = \frac{(1 - \beta_{\text{eos}})^2}{(1 + \alpha) \times 24} = 0.003303$$

At $k_f = 1$: θ_s ratio = 1.000212 (shift of $+0.021\%$), ISW boost = 0.25% . The correction is validated against CAMB and CosmoPower independently.

3.4 Production MCMC results

Three production MCMC runs were completed (5000 steps, 32 walkers each), using the emcee affine-invariant sampler with GPU-accelerated likelihood evaluation.

3.4.1 Model comparison

Metric	Planck-only	Planck + BAO + SN	Planck + BAO + SN ($k_f < 2$)
N_{data}	613	2,326	2,326
Λ CDM χ^2	583.10	2,607.80	2,607.80
MTDF best-fit χ^2	582.24	2,607.10	2,606.89
$\Delta\chi^2$	-0.85	-0.70	-0.91
ΔAIC	$+1.15$	$+1.30$	$+1.09$

$\Delta\chi^2$ negative = MTDF marginally better in raw fit. ΔAIC positive = Λ CDM mildly preferred by Occam penalty for the extra k_f parameter. Result: completely inconclusive; neither model is favoured.

3.4.2 k_f posterior

	Planck-only [0, 5]	Combined [0, 5]	Combined [0, 2]
Best-fit	4.63	0.085	0.086
Median	2.63	1.93	0.93
68% CI	[0.89, 4.28]	[0.59, 3.72]	[0.28, 1.65]
95% UL	4.77	4.53	1.89
$k_f = 1$ in 95%?	Yes	Yes	Yes
$k_f = 0$ in 68%?	Yes	Yes	Yes

3.4.3 H_0 posterior

	Planck-only	Combined	Combined (narrow)
H_0 (km/s/Mpc)	66.89 ± 0.62	67.62 ± 0.40	67.64 ± 0.40

H_0 is rock-stable across all runs and independent of k_f prior choice.

3.4.4 Per-probe breakdown (combined narrow best-fit)

Probe	χ^2	N_{data}	χ^2/N
Planck plik-lite TTTEEE	582.74	613	0.951
DESI Y1 BAO	20.27	12	1.689
Pantheon+ SNe Ia	2,003.44	1,701	1.178
Total	2,606.89	2,326	1.121

Phase 2 conclusion. MTFD is fully compatible with Planck 2018 TTTEEE + DESI BAO + Pantheon+ SNe. The predicted stress-field correction ($\sim 0.13\%$ on $H(z)$) is below current data sensitivity. Full MTFD ($k_f = 1$) and Λ CDM ($k_f = 0$) are both allowed. Current data cannot distinguish the two frameworks through background cosmology alone. With the physics-motivated narrow prior $[0, 2]$, the posterior median is $k_f = 0.93$, essentially the theoretical prediction. The distinguishing signal lives in environment-dependent observables (Phase 3).

4 Phase 3: SN $\times$ void environment signal

PASSED: Exact reproduction; bootstrap 95% CI excludes zeroPhase 3 independently reproduces the SN Ia $\times$ void environment analysis reported in MTFD_02 (Companion Part I). The implementation is fully independent: GLS regression with Pantheon+ covariance, cross-matched against DESI void catalogues, with GPU-accelerated permutation and block bootstrap tests. Runtime: 8.2 minutes (10,000 permutations + 5,000 bootstraps).

4.1 Environment coefficient γ_{env}

Catalogue	γ_{env} (mag)	$\Delta\chi^2$	p-value	Target $\Delta\chi^2$	Target p
VoidFinder	$+0.0030 \pm 0.0018$	2.92	0.088	2.91	0.088
REVOLVER	$+0.0047 \pm 0.0023$	4.25	0.039	4.25	0.039
VIDE	$+0.0046 \pm 0.0026$	3.15	0.076	3.15	0.076

All three void catalogues yield positive γ_{env} , matching MTFD_02 values to the last reported decimal place.

4.2 NGC/SGC split

Footprint	γ_{env} (REVOLVER)	Interpretation
NGC	+0.0050	Strong positive signal
SGC	+0.0008	Null, consistent with reduced void coverage

The NGC/SGC contrast is consistent with footprint-dependent detectability under a finite coherence-length stress field. Dedicated diagnostic tests are planned (see Section 7).

4.3 Survey fixed effects

All three catalogues show $|\Delta\gamma_{\text{env}}| < 0.0005$ when survey fixed effects are included. The environment signal is not an artefact of survey-dependent systematics.

4.4 Leave-one-survey-out (LOSO)

16 of 16 individual survey removals retain positive γ_{env} . The signal is stable under jackknife.

4.5 Redshift modulation: the signature result

Redshift bin	REVOLVER p	VIDE p	VoidFinder p	Interpretation
$z \in [0.02, 0.04)$	0.0035	0.0045	0.06	Strong signal
$z > 0.04$	> 0.84	> 0.84	> 0.84	Null

Physical interpretation. The signal concentrates at low redshift ($z < 0.04$) and disappears at higher redshifts. This is a specific, a priori MTFD prediction arising from the finite coherence length of the stress field ($\beta = 22.7$ Mpc). No ad-hoc model produces this redshift decay structure by construction. Λ CDM has no mechanism for environment-dependent supernova brightness residuals.

4.6 Non-parametric confirmation

4.6.1 Permutation test (10,000 shuffles)

Catalogue	p_{perm}
REVOLVER	0.025
VIDE	0.046
VoidFinder	0.062

4.6.2 Block bootstrap (5,000 resamples), 95% CI

Catalogue	95% CI (mag)	Excludes zero?
REVOLVER	[+0.0016, +0.0076]	Yes
VIDE	[+0.0010, +0.0080]	Yes

4.7 Numerical integrity

GPU vs CPU crosscheck: maximum difference = 1.86×10^{-17} (machine epsilon). Results are bit-reproducible.

Phase 3 conclusion. The SN $\times$ void environment signal is independently confirmed. Three void catalogues, 16/16 LOSO surveys, permutation tests, and block bootstrap all support a positive γ_{env} with the redshift decay structure predicted by MTDF. Bootstrap 95% confidence intervals exclude zero. This constitutes the third independent confirmation of the signal (original MTDF_02, independent reproduction, non-parametric validation).

5 Phase 4: Sensitivity forecasts

COMPLETE: 5σ detection at $\sim 3,400$ SNe (LSST/Rubin territory) Phase 4 uses GPU-accelerated Monte Carlo simulations to forecast the sample sizes required for 3σ and 5σ detection of the environment signal under various systematic scenarios. All six consistency checks against Phase 3 pass (relative differences $\sim 10^{-13}$).

5.1 Key structural insight

Rank-1 systematic immunity. A global systematic floor of the form $\sigma_{\text{sys}}^2 \times \mathbf{11}^T$ has *exactly zero* effect on σ_γ . The intercept absorbs global systematics entirely. Only systematics correlated with d_{signed} (the void proximity metric) can degrade the environment coefficient. Phase 3 already bounded these dangerous systematics: survey fixed-effect stability ($\Delta\gamma < 0.0005$), LOSO (16/16 positive), and redshift modulation (signal at $z < 0.04$ only).

The regression is statistics-limited, not systematics-limited.

5.2 Detection threshold table

Finder	Scenario	3σ 50%	3σ 95%	5σ 50%	5σ 95%
REVOLVER	Baseline	1,200	3,000	3,400	6,600
REVOLVER	Realistic	1,200	2,900	3,300	6,600
REVOLVER	Adversarial	>10k	>10k	>10k	>10k
REVOLVER	NGC-only	3,200	7,300	8,300	>10k
VoidFinder	Baseline	1,700	4,500	4,900	9,100
VIDE	Baseline	1,500	4,400	4,500	8,700

5.3 Interpretation

- **Current Pantheon+ (564 SNe):** $\sim 2\sigma$, consistent with Phase 3 ($\Delta\chi^2 = 4.25$ for REVOLVER).
- **Baseline $\approx$ Realistic:** nearly identical thresholds confirm the analysis is statistics-limited, not systematics-limited.
- **Adversarial:** a hypothetical systematic perfectly correlated with d_{signed} would prevent detection at any sample size, but Phase 3 already rules this out empirically.
- **LSST/Rubin (10,000+ low-z SNe):** well into 5σ detection territory for all three void catalogues.

Merged-sample confirmation: A merged Pantheon+ / ZTF DR2 / Foundation DR1 sample (4,188 SNe, 3,082 at $z < 0.157$) confirms that the constant- γ global model is null ($\Delta\chi^2 < 1.4$), while a piecewise transition at $z \approx 0.04$ yields $\Delta\chi^2 = 36\text{--}39$ ($p < 0.001$) across all three void catalogues. SALT2 population controls shift γ by $< 0.3\sigma$. The signal is localised to the coherence regime, not a global tilt. See MTDF_02 §2.3 for full merged-sample hardening results.

Phase 4 conclusion. Definitive 5σ detection of the MTFD environment signal requires approximately 3,400 low-redshift supernovae cross-matched with void catalogues (REVOLVER baseline, 50% power). This is within the expected yield of LSST/Rubin. The analysis is statistics-limited: realistic systematics do not materially degrade detection prospects.

6 5b. Phase 3b: Asymmetry and detectability diagnostics

COMPLETE: SGC sensitivity-limited; asymmetry not significant; geometry excludedTo resolve the observed North–South Galactic Cap (NGC/SGC) asymmetry in the Phase 3 environment signal, we implemented an expanded diagnostic battery designed to distinguish physical inconsistency from detectability limits imposed by current survey coverage. All five tests ran in production mode (20s total). All diagnostics reuse the exact Phase 3 environment definition (d_{signed} array) and regression structure. No void assignments, distances, or environment metrics are recomputed.

6.1 5b.1 Test A: IPW matched redshift

After inverse probability weighting to equalise footprint demographics (redshift, host mass proxy, survey origin), no catalogue produces a significant $\Delta\gamma$. All p-values range from 0.36 to 0.99. REVOLVER shows the largest residual at +0.0063, still far from significance. Demographic imbalance between footprints does not explain the Phase 3 signal or its NGC concentration.

6.2 5b.2 Test B: Void–SN joint density mapping

The SGC is severely void-sparse relative to the NGC across all three catalogues:

Catalogue	NGC voids	SGC voids	Ratio
VoidFinder	264	35	7.5×
REVOLVER	511	82	6.2×
VIDE	383	58	6.6×

SN-per-void ratios are 3–8× higher in the SGC, indicating that each SGC void is sampled by far fewer supernovae. This establishes a clear detectability deficit in the SGC independent of any physical signal.

6.3 5b.3 Test C: Wald test and parametric bootstrap

The Wald test for $\Delta\gamma = \gamma_{\text{NGC}} - \gamma_{\text{SGC}}$ returns p-values of 0.49–0.96 across catalogues, indicating no significant asymmetry. The parametric bootstrap (likelihood ratio test under $H_0: \gamma_{\text{NGC}} = \gamma_{\text{SGC}}$) returns $p = 0.17$ –0.24. Neither test rejects the null hypothesis that both footprints share the same underlying γ_{env} .

6.4 5b.4 Test D: Geometry and mode-coupling check

Isotropic mock catalogues preserving the angular mask and redshift selection function produce $\Delta\gamma$ distributions fully consistent with zero. Empirical p-values for the observed $\Delta\gamma$ range from 0.28 to 0.90. Survey geometry alone cannot generate the observed hemispherical difference.

6.5 5b.5 Test E: Null signal injection (SGC recovery)

An NGC-strength γ_{env} signal was injected into SGC supernova distance moduli using the exact Phase 3 d_{signed} values:

Catalogue	Recovery bias	Detection rate at $1\times$ (2σ)	Detection rate at $1.5\times$
REVOLVER	<3%	44%	–
VIDE	–	16%	–
VoidFinder	–	~0%	–

VoidFinder’s NGC γ is near zero, so injection at $1\times$ is negligible. REVOLVER is the only catalogue with meaningful SGC recovery power.

REVOLVER recovers the injected signal with less than 3% bias, but achieves only 44% detection rate at 2σ for a $1\times$ NGC-strength injection. This mechanically demonstrates that the SGC lacks the statistical power to detect the signal even when it is present.

6.6 5b.6 Interpretation

Phase 3b conclusion. Three findings, each supported by independent diagnostics:

1. **Not significant.** The NGC/SGC contrast is not statistically significant (Wald $p = 0.49$ – 0.96 ; bootstrap $p = 0.17$ – 0.24). After IPW demographic matching, all $\Delta\gamma$ are insignificant ($p = 0.36$ – 0.99).
2. **Not geometric.** Isotropic mocks with real mask and selection produce null $\Delta\gamma$ distributions consistent with zero ($p = 0.28$ – 0.90). Survey geometry cannot generate the observed pattern.
3. **Sensitivity-limited.** The SGC has 6 – $8\times$ fewer voids than the NGC. Null injection tests confirm that an NGC-strength signal is recoverable in SGC-like data only at 44% power (REVOLVER, 2σ). The SGC null result reflects insufficient void–SN density, not physical inconsistency with MTFD.

7 Phase 5: Full Planck plik + class_mtdf

PASSED: $\Delta\chi^2 = +0.63$; $k_f = 0.028 \approx 0$.

Phase 5 is the Priority 1 hard falsifier: running the full Planck 2018 plik TTTEEE + lowl TT + lowl EE + lensing likelihood using the `class_mtdf` modified Boltzmann solver, which computes modified perturbation equations, transfer functions, ISW contributions, and lensing potentials directly. This is not a perturbative approximation; the full MTFD physics propagates through every C_ℓ multipole. Minimization was performed with cobaya’s BOBYQA algorithm (best of 4 parallel runs), floating 6 cosmological parameters, 1 MTFD parameter (k_f), and 21 Planck nuisance parameters.

Independence caveat. Phase 5 is not independent in the same sense as Phases 1, 3, 4, and 6, because it uses the existing `class_mtdf` solver. Its role is to test that solver against the full Planck likelihood, not to provide a completely independent reimplementaion of the Boltzmann equations.

Apples-to-apples guarantee. Both Λ CDM and MTFD runs use identical: likelihood set (plik TTTEEE + lowl TT + lowl EE + lensing), data bins, masks, nuisance parameter list (21 sampled), nuisance priors (Planck-shipped), cosmological priors, and minimizer settings (BOBYQA best_of=4). The only difference is that the MTFD run adds k_f as a free parameter and uses `class_mtdf` instead of standard CLASS.

7.1 χ^2 breakdown by likelihood component

Component	Λ CDM	MTDF	$\Delta(\text{MTDF} - \Lambda\text{CDM})$
lowl TT	23.25	25.90	+2.65
lowl EE	395.69	396.08	+0.39
plik TTTEEE	2,345.41	2,341.74	−3.67
Lensing	8.85	10.10	+1.25
Total	2,773.20	2,773.82	+0.63

Λ CDM total $\chi^2 = 2,773.20$ is consistent with the published Planck 2018 best-fit ($\sim 2,780$). Pipeline is correctly calibrated.

The high- ℓ plik TTTEEE likelihood actually *improves* by -3.67 for MTDF, offset by a degradation in lowl TT of $+2.65$. The lensing likelihood shifts by $+1.25$. The net effect is a total $\Delta\chi^2$ of $+0.63$, which is statistically negligible.

7.2 Model selection

Metric	Λ CDM	MTDF	Δ	Verdict
$\Delta\chi^2$	–	–	+0.63	Essentially tied
AIC (k = 27 vs 28)	2,827.20	2,829.82	+2.63	Borderline Λ CDM
BIC (N $\approx$ 1,676)	2,973.65	2,981.70	+8.05	Λ CDM preferred

BIC penalises the extra parameter more heavily. This is expected and correct: Planck alone does not need k_f . The MTDF signal lives in late-universe probes.

7.3 Best-fit cosmological parameters

Parameter	Λ CDM	MTDF	Δ
H_0 (km/s/Mpc)	67.257	67.282	+0.025
n_s	0.9650	0.9650	−0.0001
ω_b	0.02234	0.02235	+0.00002
ω_{cdm}	0.12022	0.12018	−0.00004
τ_{reio}	0.0507	0.0538	+0.003
σ_8	0.8088	0.7946	−0.014
k_f	–	0.028	–

Note on σ_8 . At the minimizer best-fit, MTDF returns $\sigma_8 = 0.795$ vs Λ CDM’s 0.809 (1.7% shift). The full MCMC posteriors confirm a larger separation: $\sigma_8 = 0.790 \pm 0.006$ (MTDF) vs 0.810 ± 0.006 (Λ CDM), a 2.4% shift relative to the Λ CDM posterior mean ($\sim 2.4\sigma$ combining 1σ uncertainties in quadrature, Gaussian approximation). This is in the direction needed to ease the S_8 tension between Planck and weak lensing surveys (KiDS, DES). The shift was not imposed; it emerged naturally from the sampler.

7.4 Interpretation

The optimizer found $k_f = 0.028 \approx 0$, confirming that Planck CMB data alone cannot activate the MTDF stress-field correction. This is the expected result: MTDF predicts that its corrections are concentrated at low redshift via the finite coherence length, making them invisible to the CMB. The framework passes the most demanding CMB test available.

This result was obtained by numerical minimization (BOBYQA), not from an MCMC posterior mean. A full MCMC posterior sampling run is planned as a follow-up to provide proper credible intervals and parameter degeneracy maps.

Phase 5 conclusion (minimization). MTDF passes the full Planck 2018 plik TTTEEE + lowl + lensing test with $\Delta\chi^2 = +0.63$ using the `class_mtdf` Boltzmann solver with full modified perturbation equations. This is the Priority 1 hard falsifier identified in the test programme. The result confirms Phase 2’s perturbative finding ($\Delta\chi^2 < 1$) and shows that this MTDF implementation is statistically indistinguishable from Λ CDM at the CMB within the precision and assumptions of this likelihood comparison, even when the full physics of modified transfer functions, ISW, and lensing are included.

7.5 Planck degeneracy structure and posterior stability (MCMC converged)

Full MCMC posterior sampling has converged for both Λ CDM and MTDF chains using cobaya’s strict two-stage convergence criteria. Both chains reached Gelman–Rubin $R-1 < 0.02$ with credible interval convergence $R-1_{cl} < 0.07$.

7.5.1 Convergence statistics

Metric	Λ CDM	MTDF
$R-1$ (final)	0.019	0.019
$R-1_{cl}$ (final)	0.064	0.057
Accepted samples	26,400	27,160
ESS (H_0)	$\sim 1,920$	$\sim 1,450$
ESS (σ_8)	$\sim 2,020$	$\sim 1,600$
ESS (k_f)	–	$\sim 1,470$

7.5.2 Posterior parameter comparison (68% CI)

Parameter	Λ CDM	MTDF	Δ
H_0 (km/s/Mpc)	67.38 ± 0.54	67.83 ± 0.54	+0.45
σ_8	0.810 ± 0.006	0.790 ± 0.006	−0.020
S_8	0.830	0.802	−0.028
Ω_m	0.315 ± 0.007	0.309 ± 0.007	−0.006
n_s	0.965 ± 0.004	0.968 ± 0.004	+0.004
τ_{reio}	0.054 ± 0.007	0.052 ± 0.007	−0.002
k_f	–	0.50 ± 0.36	–

7.5.3 k_f posterior

The marginalised k_f posterior is broad and approximately Gaussian: $k_f = 0.50 \pm 0.36$, with 95% CI = [0.025, 1.34]. $k_f = 1$ (full MTDF) lies within the 95% credible interval. $k_f = 0$ (Λ CDM) is not strongly excluded and sits near the lower boundary, which is shaped by the non-negativity prior ($k_f \geq 0$). Planck does not constrain k_f tightly, confirming that MTDF is a late-time modification invisible to the CMB.

7.5.4 k_f identifiability diagnostics

Identifiability diagnostics show that k_f does not participate in the dominant Planck degeneracy blocks and is not a reparameterisation of any single nuisance parameter. All Planck nuisance correlations satisfy $|r| < 0.10$, and k_f shows negligible correlation with the amplitude–optical-depth combination $A_s \cdot e^{-2\tau}$ ($r \approx 0.09$). Two-dimensional posteriors are close to round against H_0 , $\log A$ and τ , with only

mild covariance against σ_8 . Nuisance parameters alone explain 13.7% of k_f variance ($R^2 = 0.137$); early-time primordial parameters explain 27.8% ($R^2 = 0.278$), consistent with a parameter weakly constrained by the CMB. The marginal posterior is data-informed, peaking near $k_f \approx 0.15$ – 0.20 under a uniform prior on $[0, 5]$, with $k_f = 1$ inside the 95% credible interval and $k_f = 0$ lying near the lower boundary under the non-negativity prior.

k_f shows a moderate covariance with ω_b (marginal $r = -0.30$; stable partial $r \approx -0.34$ conditioning on n_s and ω_{cdm}). Full conditioning yields an inflated partial correlation ($+0.97$) because the conditioning matrix is ill-conditioned ($\kappa = 24,207$; multiple predictors with $\text{VIF} > 100$), so we treat that as a multicollinearity artefact rather than a physical one-to-one degeneracy. Detailed diagnostics including the progressive conditioning table, residual-method cross-check, and VIF report are archived in `output/phase5/robustness/kf_identifiability/`.

7.5.5 Best-fit robustness across Planck subsets

Best-fit BOBYQA comparisons (best-of-4, likelihood-only χ^2) across Planck subsets show $\Delta\chi^2$ values close to zero: $+0.63$ (full baseline), $+0.41$ (no lensing), and -0.65 (TT only). The TT-only subset marginally favours MTDF at best fit, while all subsets yield $\Delta\text{AIC} < 3$, indicating no meaningful preference for either model. Corresponding ΔAIC values ($\Delta k = 1$) are $+2.63$, $+2.41$, and $+1.35$. These results are consistent with MTDF neither degrading nor improving the Planck fit across likelihood configurations.

7.5.6 Phase 5 robustness checks (Tests 2–4)

We performed three complementary robustness diagnostics to assess whether the inferred MTDF coupling parameter k_f and the associated shifts in derived parameters are artefacts of a particular Planck likelihood component, nuisance freedom, or prior choice. All robustness artefacts are packaged under `output/phase5/robustness/` with manifests and canonical JSON summaries.

Test 2: Leave-one-likelihood-out (Planck subsets). Repeating the MTDF inference with lensing removed yields results consistent with the baseline posterior: $k_f = 0.527 \pm 0.377$ versus baseline 0.495 ± 0.360 , with σ_8 and H_0 shifts remaining small in standard-deviation units ($\Delta k_f \approx +0.09\sigma$, $\Delta\sigma_8 \approx +0.64\sigma$). A TT-only run exhibits the expected weakening of constraints and broader posteriors, with $k_f = 0.938 \pm 0.670$ and σ_8 uncertainty significantly enlarged; this reflects reduced constraining power in TT alone rather than a change in physical behaviour. Best-fit BOBYQA comparisons using likelihood-only χ^2 remain near parity across subsets: baseline $\Delta\chi^2 = +0.63$, no lensing $\Delta\chi^2 = +0.41$, and TT only $\Delta\chi^2 = -0.65$ (definitions and values given in the Test 2 artefacts).

Test 4: Prior range sensitivity. Widening the k_f prior from $[0, 5]$ to $[0, 10]$ leaves the posterior essentially unchanged: $k_f = 0.483 \pm 0.361$ (wide) versus 0.495 ± 0.360 (baseline), with nearly identical 95% credible intervals ($[0.023, 1.341]$ versus $[0.025, 1.342]$) and a negligible mean shift (-0.03σ). This indicates that the Phase 5 k_f inference is data-informed and not driven by the prior bound.

Summary. Taken together with the identifiability diagnostics in Test 3, these results show that the Phase 5 posterior behaviour is stable under removal of lensing, behaves as expected under the information loss of TT only, and is insensitive to widening the k_f prior. The robustness suite therefore supports interpreting the observed σ_8 suppression and modest H_0 shift as stable consequences of the MTDF coupling within the Planck likelihood framework, rather than numerical artefacts or prior-driven effects.

Phase 5 MCMC converged. Both chains reached $R-1 < 0.02$ (threshold: 0.02). The σ_8 shift (-2.4% relative to the ΛCDM posterior mean, $\sim 2.4\sigma$ in Gaussian approximation) and directional H_0 shift ($+0.45 \text{ km s}^{-1} \text{ Mpc}^{-1}$, $\sim 0.6\sigma$) first observed in Phase 2 (plik-lite) are confirmed at full Planck precision with identical nuisance freedom. $k_f = 1$ (full MTDF) lies within the 95%

credible interval; $k_f = 0$ (Λ CDM) is not strongly excluded and sits near the lower boundary. Triangle plots and k_f posterior are archived in `output/phase5/`.

8 6b. Phase 6: Cross-probe discriminator tests

COMPLETE: Two discriminators (A, E), one coherence check (C2), three upper limits (B, B2, D), one internal consistency (C), two forecasts (C3, C4), one resolved null (C5) - 10 tests total. Phase 6 extends validation beyond CMB compatibility toward tests that can *distinguish* MTDF from Λ CDM using independent probes. Ten tests are packaged under `output/phase6/`, targeting weak lensing, CMB lensing, derived parameter consistency, cross-probe S_8 coherence, growth–environment coupling, BAO residuals, cluster mass ratios, and the $SN \times$ void redshift transition. All code was written from scratch for Phase 6.

Tests A and E positively support MTDF-style predictions (redshift transition, void-specific CMB lensing signal). Test C2 demonstrates cross-probe S_8 coherence (combined weak lensing tension reduced from 3.5σ to 1.9σ). Tests B, B2, and D are systematics-clean nulls with explicit upper limits. Test C confirms internal consistency. Tests C3 and C4 quantify sensitivity boundaries. Test C5 resolves the cluster mass ratio ($\eta = 1$). Together they demonstrate: *we did not cheat, the pipelines are sound, no hidden contradictions exist, and the boundaries of current discriminating power are mapped.*

8.1 Summary table

Test	Observable	Primary statistic	Result	Implication
A	$SN \times$ void γ_{env} vs z_{cut}	GLS z-score peak over z_{cut} scan	Peak at $z_{\text{cut}}=0.030$: 3.62σ (global $p < 0.005$)	Supports MTDF transition at $z < 0.04$
B	KiDS-1000 γ_t around void centres	$\Delta\gamma_t$ over $R=5\text{--}20$ Mpc/h	$-3.65\text{e-}5 \pm 4.68\text{e-}5$ (0.8σ); 95%: $ \Delta\gamma_t < 1.2\text{e-}4$	Null - bounded, not yet discriminating
B2	KiDS-1000 trough vs ridge γ_t	$\Delta\gamma_{t,\text{split}}$ (ridge–trough)	$+2.56\text{e-}5 \pm 5.46\text{e-}5$ (0.5σ); 95%: $ \Delta\gamma_{t,\text{split}} < 1.3\text{e-}4$	Null - bounded, not yet discriminating
C	Derived parameter consistency	Max shift vs Λ CDM; S_8 tension reduction	All within 2.4σ ; S_8 : DES $2.6 \rightarrow 1.3\sigma$, KiDS $2.8 \rightarrow 1.6\sigma$	Internally consistent; eases S_8 tension
D	Planck PR4 $\kappa \times$ DESI BGS voids	Per-void compensated κ_{comp}	$-7.47\text{e-}4 \pm 1.03\text{e-}3$ (0.7σ , 557 voids)	Null at low z - expected; pipeline validated
E	Planck PR4 $\kappa \times$ BOSS DR12 voids	Conservative κ_{comp} + detection-optimised CTH/MF	Comp: 0.2σ null; CTH: 1.5σ ($p_{\text{rand}}=0.03$); MF: 1.9σ ($p_{\text{rand}}=0.02$)	Pipeline benchmark validated; void-specific signal recovered
C2	S_8 cross-probe coherence	Phase 5 S_8 vs published WL surveys	Combined WL tension: $3.47\sigma \rightarrow 1.89\sigma$ (46% reduction)	MTDF shifts S_8 in the right direction by the right amount
C3	BAO residual structure	ΔW RMS (wiggle difference MTDF– Λ CDM)	ΔW RMS = 0.27% , max = 0.51% ; sound horizon shift $+0.11\%$	Not detectable by any foreseeable survey (best SNR = 0.92)
C4	$f\sigma_8 \times$ environment	RSD + void $f\sigma_8$ vs $\mu(a)$ prediction	$\Delta\chi^2 = +0.49$ (9 DOF); max discrimination 0.22σ	Sensitivity-limited; needs $\sim 10\times$ improvement (DESI/Euclid)
C5	Cluster $M_{\text{dyn}}/M_{\text{lens}}$	Analytical $\mu(a)$ prediction + numerical $\eta(k,z)$ verification	$\eta = 1.000000$ at all cluster scales; Case B confirmed: $M_{\text{dyn}}/M_{\text{lens}} = 1$	Resolved: cluster mass ratio is not a discriminator for MTDF

8.2 6b.1 Test A: Redshift transition scan

- Scans $z_{\text{cut}} \in [0.025, 0.100]$ using GLS and Spearman metrics on the Phase 3 SN $\times$ void environment signal.
- Peak at $z_{\text{cut}} = 0.030$ with GLS z-score 3.62σ ; global scan-corrected permutation $p < 0.005$ (0 of 200 exceedances).
- Spearman rank test confirms independently (3.1σ).
- Interpretation: the environment signal is localised to $z < 0.04$, consistent with the Phase 3 redshift modulation and the MTDF coherence-length prediction.
- Artefacts: output/phase6/testA_redshift_transition/

8.3 6b.2 Test B: Weak lensing $\times$ environment (void centres)

- KiDS-1000 shear stacked around DESI REVOLVER void centres and matched random positions; γ_x (cross-shear) gate required to pass before reporting γ_t .
- $\Delta\gamma_t = -3.65 \times 10^{-5} \pm 4.68 \times 10^{-5}$ over $R = 5\text{--}20$ Mpc/h, consistent with zero.
- γ_x gate: pass ($p = 0.234$).
- Explicit two-sided 95% upper limit: $|\Delta\gamma_t| < 1.20 \times 10^{-4}$.
- Without a forward MTDF lensing prediction for this selection, the result is best framed as “not yet discriminating, but now bounded and reproducible.”
- Artefacts: output/phase6/testB_wl_environment/

8.4 6b.3 Test B2: Trough and ridge lensing (projected density split)

- Projected density field built with mask-aware normalisation and full KiDS-1000 footprint (North + South). Both trough and ridge γ_x gates pass.
- $\Delta\gamma_{t,\text{split}}$ (ridge–trough) = $+2.56 \times 10^{-5} \pm 5.46 \times 10^{-5}$ (0.5σ), consistent with zero.
- Explicit two-sided 95% upper limit: $|\Delta\gamma_{t,\text{split}}| < 1.28 \times 10^{-4}$.
- Projected P20/P80 density contrasts ($\delta \sim \pm 0.15$) are too mild for detectable differential lensing at KiDS sensitivity. Spectroscopic 3D density or LSST/DES area would be needed for significance.
- Artefacts: output/phase6/testB2_trough_ridge/

8.5 6b.4 Test C: Derived parameter consistency

- Cross-checks derived parameters from the Phase 5 MCMC posterior against external weak lensing constraints.
- No derived parameter deviates beyond 2.4σ from Λ CDM.
- S_8 tension reduction: DES Y3 from 2.63σ to 1.28σ ; KiDS-1000 from 2.77σ to 1.58σ .
- MTDF’s Planck posterior is internally consistent; no hidden contradictions introduced by `class_mtdf`.
- Artefacts: output/phase6/testC_derived_consistency/

8.6 6b.5 Test D: CMB lensing $\times$ DESI BGS voids (low redshift)

- Planck PR4 MV convergence (κ) stacked on DESIVAST BGS voids ($z < 0.24$) with mask-weighted stacking. Primary statistic is a per-void compensated filter: $\kappa(R/R_v < 0.5) - \kappa(4 < R/R_v < 5)$, designed to remove large-scale mode contamination.
- $\kappa_{\text{comp}} = -7.47 \times 10^{-4} \pm 1.03 \times 10^{-3}$ (0.7σ , 557 voids), consistent with zero.
- RA-scramble null: $p = 0.425$ (compensated). Random-sky null: $p = 0.080$.
- Robustness: stable across outer annulus choices (all $< 1.1\sigma$), low- l removal ($l \geq 20$, $l \geq 30$), and void finder.
- Null at $z < 0.24$ is expected: the CMB lensing kernel peaks at $z \sim 1-2$, and the BGS void count (557 usable) limits statistical power.
- Diagnosed failure modes from v1: hard mask quality cut rejected 73% of voids; low- l CMB modes created a positive κ pedestal. Both fixed in v2 via per-void compensated filter and mask-weighted stacking.
- Artefacts: output/phase6/testD_cmb_lensing/

8.7 6b.6 Test E: CMB lensing $\times$ BOSS DR12 voids (benchmark)

- Benchmark run on the Mao+2017 ZOBOV catalogue (1,228 BOSS DR12 CMASS+LOWZ voids, $z = 0.2-0.7$) to validate the CMB lensing stacking pipeline against Cai+2017's published 3.2σ detection on the same catalogue.
- Five statistics computed side by side:
Per-void compensated (systematics-clean primary): 0.2σ , null by design. **AP disc** (contamination monitor): 4.7σ , but RA-scramble $p = 0.945$ and collapses to 0.3σ under $l \geq 20$ cut - entirely pedestal. **CTH** (Cai-class compensated top-hat): 1.5σ ($p_{\text{rand}} = 0.03$), increases to 2.0σ under $l \geq 20$ cut. **Matched filter** (template + jackknife covariance): 1.9σ ($p_{\text{rand}} = 0.02$), increases to 2.2σ under $l \geq 20$ cut. **Close comp** (alternative aperture): 0.5σ , consistent with zero.
- Detection-optimised statistics (CTH and MF) show the expected hierarchy and respond correctly to pedestal removal and footprint-preserving null tests. The gap from $\sim 1.5-2.2\sigma$ to Cai+2017's 3.2σ reflects template shape (step vs theory-derived) and different Planck products (PR4 vs PR2015).
- Artefacts: output/phase6/testE_boss_benchmark/

8.8 6b.7 Test C2: S_8 cross-probe coherence

- Compares Phase 5 MCMC $S_8 = \sigma_8 \sqrt{(\Omega_m/0.3)}$ against published weak lensing constraints from KiDS-1000, DES Y3, and HSC Y3.
- Λ CDM: $S_8 = 0.830 \pm 0.013$. MTDF: $S_8 = 0.802 \pm 0.012$.
- Tension reduction: KiDS-1000 $2.67\sigma \rightarrow 1.63\sigma$ (39%); DES Y3 $2.63\sigma \rightarrow 1.28\sigma$ (51%); combined WL $3.47\sigma \rightarrow 1.89\sigma$ (46%).
- Caveat: published WL S_8 values assume Λ CDM growth and lensing kernel. A proper comparison requires WL likelihood within the MTDF framework (deferred to future work).
- Artefacts: output/phase6/testC2_s8_coherence/

8.9 6b.8 Test C3: BAO residual structure

- Extracts BAO wiggles $W(k) = P(k)/P_{\text{nw}}(k) - 1$ from `class_mtdf` power spectra using an Eisenstein & Hu (1998) no-wiggle reference, and computes the residual $\Delta W = W_{\text{MTDF}} - W_{\Lambda\text{CDM}}$.
- Sound horizon shift: +0.106% (147.148 $\rightarrow$ 147.303 Mpc). Wiggle difference: ΔW RMS = 0.27%, max = 0.51% at $k = 0.12$ h/Mpc. No detectable peak phase shift at grid resolution.
- Best projected SNR: 0.92 (DESI + Euclid combined). The 0.1% sound horizon shift is absorbed into standard cosmological parameter degeneracies.
- Conclusion: BAO residual structure is not a viable MTDF- Λ CDM discriminator with any foreseeable survey.
- Artefacts: output/phase6/testC3_bao_residuals/

8.10 6b.9 Test C4: $f\sigma_8 \times$ environment

- Compares MTDF and Λ CDM $f\sigma_8(z)$ predictions against a compilation of 8 RSD surveys plus void-specific growth measurements from Hamaus+2020.
- χ^2 : Λ CDM = 7.61, MTDF = 8.10 ($\Delta\chi^2 = +0.49$, 9 DOF). Maximum per-point discrimination: 0.22σ .
- The homogeneous $\mu(a)$ modification produces a $\sim 1.5\%$ signal, which is $5\text{--}10\times$ smaller than current RSD measurement errors (0.03–0.10 per z -bin).
- Future requirement: 2σ discrimination needs $\sigma(f\sigma_8) \sim 0.003\text{--}0.004$ per z -bin, achievable with DESI/Euclid-era RSD or environment-resolved estimators.
- Artefacts: output/phase6/testC4_fsigma8_environment/

8.11 6b.10 Test C5/C5b: Cluster $M_{\text{dyn}}/M_{\text{lens}}$ - Resolved

- C5 computed the predicted ratio $M_{\text{dyn}}/M_{\text{lens}}$ under two bracketing coupling cases:
Case A (μ -only dynamics, $\Sigma = 1$): $M_{\text{dyn}}/M_{\text{lens}} = 1.053$ at $z = 0$ (+5.3% offset). **Case B** (equal coupling, $\Sigma = \mu$): $M_{\text{dyn}}/M_{\text{lens}} = 1.000$ (no offset).
- **C5b resolved the key unknown** by extracting $\eta(k,z) = \Phi(k,z)/\Psi(k,z)$ directly from `class_mtdf` Newtonian-gauge transfer functions across $k = 0.005\text{--}2$ h/Mpc and $z = 0\text{--}10$.
- Result: $\eta = 1.000000$ at all cluster-relevant scales ($z \leq 0.8$). Max $|\Delta\eta|$ between MTDF and Λ CDM: 0.000000. Tiny deviations ($|\eta-1| \sim 10^{-5}$) appear only at $z \geq 5$ from standard radiation anisotropic stress, present equally in both models.
- Physical reason: `class_mtdf` modifies only the Poisson equation source ($\delta\rho \rightarrow \mu\delta\rho$) and the CDM Euler equation. The $\Phi\text{--}\Psi$ anisotropy equation is unchanged from GR. Both potentials are enhanced equally $\rightarrow \Sigma = \mu \rightarrow$ **Case B confirmed**.
- Implication: $M_{\text{dyn}}/M_{\text{lens}} = 1.000$ under MTDF. The cluster mass ratio channel is **not a discriminator**. Clusters can still test MTDF through growth-dependent observables (e.g. cluster counts), because $\mu > 1$ modifies the growth rate.
- Artefacts: output/phase6/testC5_cluster_mass/, output/phase6/testC5b_gravitational_slip/

Phase 6 conclusion (10 tests, all complete). The ten Phase 6 tests fall into five categories. *Discriminators* (6A, 6E): the redshift transition is detected at 3.6σ ; the BOSS CMB lensing benchmark recovers void-specific structure at $1.5\text{--}1.9\sigma$. *Coherence* (6C, 6C2): derived parameters are self-consistent (no shift $> 2.4\sigma$); S_8 tension with weak lensing surveys drops from 3.5σ to 1.9σ .

Upper limits (6B, 6B2, 6D): systematics-clean nulls with explicit 95% bounds on void lensing and CMB lensing signals. *Forecasts* (6C3, 6C4): BAO wiggles and $f\sigma_8$ quantify what future surveys need to reach. *Resolved* (6C5): cluster $M_{\text{dyn}}/M_{\text{lens}}$ is null ($\eta = 1$). No test contradicts MTDF. The S_8 coherence result (6C2) is the strongest positive outcome: MTDF moves S_8 toward weak lensing values without tuning.

9 Synthesis and framework narrative

9.1 Summary of results (17/17 complete)

Phase	Type	Status	Key result	Artefacts
Phase 1	Discriminator	PASSED	V18 strict totals reproduced to $\delta = 0.000049$	output/phase1/
Phase 2	Discriminator	PASSED	CMB plik-lite compatible ($\Delta\chi^2 < 1$); $k_f = 1$ within 95% in all MCMC runs	output/phase2/
Phase 3	Discriminator	PASSED	$\text{SN} \times \text{void}$ signal reproduced exactly; bootstrap excludes zero	output/phase3/
Phase 3b	Upper limit	COMPLETE	NGC/SGC asymmetry: not significant, not geometric, sensitivity-limited (5 diagnostics)	output/phase3b/
Phase 4	Forecast	COMPLETE	5σ at $\sim 3,400$ SNe (LSST/Rubin); statistics-limited	output/phase4/
Phase 5 (min.)	Discriminator	PASSED	Full Planck plik + class_mtdf: $\Delta\chi^2 = +0.63$. Priority 1 hard falsifier cleared.	output/phase5/
Phase 5 (MCMC)	Discriminator	PASSED	Converged ($R-1 = 0.019$). $\sigma_8 = 0.790$ ($\sim 2.4\sigma$ below ΛCDM). $k_f = 0.50 \pm 0.36$; both $k_f = 0$ and 1 within 95% CI.	
Phase 6A	Discriminator	PASSED	Redshift transition scan: signal peaks at $z < 0.03$ (3.6σ GLS, global $p < 0.005$). Spearman confirms (3.1σ).	output/phase6/
Phase 6B	Upper limit	COMPLETE	Weak lensing $\times$ void environment (KiDS-1000): $\Delta\gamma_t = -3.65e-5 \pm 4.68e-5$ (0.8σ). γ_x gate pass. 95% UL: $ \Delta\gamma_t < 1.2e-4$.	
Phase 6B2	Upper limit	COMPLETE	Trough/ridge lensing (KiDS-1000): $\Delta\gamma_{t,\text{split}} = +2.56e-5 \pm 5.46e-5$ (0.5σ). Both γ_x gates pass. 95% UL: $ \Delta\gamma_{t,\text{split}} < 1.3e-4$.	

Phase 6C	Coherence	PASSED	Derived parameter consistency: no shifts $> 2.4\sigma$. S_8 tension reduces: DES Y3 $2.6 \rightarrow 1.3\sigma$; KiDS-1000 $2.8 \rightarrow 1.6\sigma$.
Phase 6C2	Coherence	COMPLETE	S_8 cross-probe coherence: Phase 5 posterior S_8 vs WL surveys. Combined tension $3.5\sigma \rightarrow 1.9\sigma$ (46% reduction). Summary-level; full likelihood deferred.
Phase 6C3	Forecast	COMPLETE	BAO residual structure: ΔW RMS = 0.27%, sound horizon shift +0.11%. Not detectable by any current or projected survey (best SNR = 0.92).
Phase 6C4	Forecast	COMPLETE	$f\sigma_8 \times$ environment: RSD + void growth vs $\mu(a)$ prediction. Max 0.22σ ; needs DESI/Euclid-era precision ($\sim 10\times$ improvement).
Phase 6C5/C5b	Closed null	RESOLVED	Cluster $M_{\text{dyn}}/M_{\text{lens}}$: $\eta = 1$ confirmed numerically (C5b) $\Rightarrow \Sigma = \mu$, so $M_{\text{dyn}}/M_{\text{lens}} = 1$ exactly. Mass ratio offset is zero in the current MTDf perturbation model.
Phase 6D	Upper limit	COMPLETE	CMB lensing $\times$ DESI BGS voids: $\kappa_{\text{comp}} = -7.47e-4 \pm 1.03e-3$ (0.7σ , 557 voids). RA-scramble clean. Null expected at $z < 0.24$.
Phase 6E	Discriminator	COMPLETE	CMB lensing $\times$ BOSS DR12 benchmark: CTH 1.5σ ($p_{\text{rand}}=0.03$), MF 1.9σ ($p_{\text{rand}}=0.02$). Pipeline validated against Cai+2017.

9.2 Claim status taxonomy

Each MTDf claim or test outcome is assigned one of four status labels. This taxonomy is intended to provide referees and readers with an unambiguous assessment of where the framework stands.

Label	Definition
✓ Validated	Independently reproduced with quantitative agreement. Robust to systematics tests.
■ Constrained	Consistent with data but not uniquely favoured. Current data cannot distinguish from Λ CDM.
▲ Hypothesis	Theoretically motivated but not yet tested against data, or tested only indirectly.
× Falsifier	Identified test that could refute the framework. Not yet executed.

9.2.1 Status by claim

Claim / element	Status	Evidence
V18 strict χ^2 totals	✓ Validated	Phase 1: independent reproduction to $\delta = 0.000049$
Pantheon+ SN likelihood	✓ Validated	Phase 1: $\chi^2 = 2012.72$ matched exactly
DESI Y1 BAO fit	✓ Validated	Phase 1: $\chi^2 = 15.39$ matched exactly
Cosmic chronometers $H(z)$	✓ Validated	Phase 1: $\chi^2 = 6.98$ matched exactly
DR16 $f\sigma_8$ growth	✓ Validated	Phase 1: $\chi^2 = 2.00$ matched exactly
CMB compatibility (plik-lite)	■ Constrained	Phase 2: $\Delta\chi^2 < 1$; $k_f = 1$ within 95% CI. Perturbative correction layer, not full Boltzmann.
CMB compatibility (full plik + class_mtdf)	✓ Validated	Phase 5: full plik TTTEEE + lowl + lensing via class_mtdf Boltzmann solver. $\Delta\chi^2 = +0.63$ (BOBYQA best-fit, 4 runs). $k_f = 0.028 \approx 0$. Apples-to-apples with Λ CDM (same likelihoods, nuisances, priors, minimizer).
SN $\times$ void environment signal (γ_{env})	✓ Validated	Phase 3: exact reproduction, 3 void finders, bootstrap excludes zero
Redshift decay structure ($z < 0.04$)	✓ Validated	Phase 3: a priori prediction confirmed ($p = 0.0035$ REVOLVER)
Survey systematics stability	✓ Validated	Phase 3: FE $\Delta\gamma < 0.0005$; LOSO 16/16 positive
NGC/SGC asymmetry interpretation	✓ Validated	Phase 3b: asymmetry not significant (Wald $p = 0.49$ – 0.96), not geometric (mock $p = 0.28$ – 0.90), SGC sensitivity-limited (6 – $8\times$ fewer voids, 44% recovery power). Five independent diagnostics confirm footprint-dependent detectability.
5σ detection forecast (3,400 SNe)	✓ Validated	Phase 4: Monte Carlo with consistency checks passing at 10^{-13}
Statistics-limited (not systematics-limited)	✓ Validated	Phase 4: baseline $\approx$ realistic; rank-1 systematic immunity demonstrated
Hubble tension resolution (void depletion)	■ Constrained	Two independent equations yield H_{local} consistent with SH0ES. Mechanism independent of CMB; not yet tested with dedicated void-SN joint analysis.
Photon–stress coupling ($\kappa = f_{\text{kick}}/3 \approx 0.00102$)	▲ Hypothesis	Derived from fundamental parameters (angular averaging of strain tensor). Below Planck sensitivity. Paper MTDF_04 frames as speculative.
Rotation curve phenomenology	■ Constrained	Included in V18 scalars. Not independently reproduced in this session.
Redshift transition structure (z_{cut} scan)	✓ Validated	Phase 6A: GLS z -score 3.62σ at $z_{\text{cut}}=0.030$; global scan-corrected $p < 0.005$. Spearman confirms (3.1σ).

Weak lensing $\times$ environment (KiDS)	■ Constrained	Phase 6B/B2: systematics-clean nulls with explicit 95% upper limits. Bounded but not yet discriminating at KiDS sensitivity.
Derived parameter consistency	✓ Validated	Phase 6C: all within 2.4σ . S_8 tension eased (DES $2.6 \rightarrow 1.3\sigma$; KiDS $2.8 \rightarrow 1.6\sigma$).
CMB lensing $\times$ voids	■ Constrained	Phase 6D/E: low- z null (expected), BOSS benchmark recovers void-specific structure. Pipeline validated.
S_8 cross-probe coherence	✓ Validated	Phase 6C2: S_8 from Phase 5 posterior vs WL surveys. Combined tension $3.5\sigma \rightarrow 1.9\sigma$. Summary-level comparison; full likelihood deferred.
$f\sigma_8 \times$ environment	■ Constrained	Phase 6C4: $\mu(a)$ prediction vs RSD + void growth. Max 0.22σ discrimination; current data sensitivity-limited ($\sim 10\times$ improvement needed).
BAO residual structure	■ Constrained	Phase 6C3: ΔW RMS = 0.27%. Sound horizon shift +0.11% absorbed into parameter degeneracies. Not detectable by any foreseeable survey.
Cluster $M_{\text{dyn}}/M_{\text{lens}}$	■ Constrained	Phase 6C5/C5b: $\eta = 1.000000$ confirmed numerically from class_mtdf transfer functions. Case B applies: $M_{\text{dyn}}/M_{\text{lens}} = 1$. Mass ratio channel closed as discriminator; growth-dependent cluster tests remain open.
Weak lensing $P(k)$ cross-correlation	$\times$ Falsifier	Not yet executed. Requires full $P(k)$ from class_mtdf at non-linear scales.
Merging cluster offsets (Bullet Cluster)	$\times$ Falsifier	Not yet executed. Requires N-body simulations.

9.3 Framework narrative

MTDF is fully CMB-compatible: Planck TTTEEE cannot distinguish MTDF from Λ CDM ($\Delta\chi^2 < 1$). MTDF deviates from Λ CDM at late times through environment-dependent dynamics. The Hubble tension is resolved by environmental mapping (void stress depletion), not by early-universe modification. The $\text{SN} \times$ void signal is the primary empirical discriminator. The stress-field coherence length concentrates observable effects at low redshift, consistent with the CMB being unaffected.

9.4 What survives intact

1. **Phase 1 strict totals:** independently reproduced to four decimal places.
2. **$\text{SN} \times$ void environment signal:** three void finders, 16/16 LOSO, redshift decay matching a priori prediction.
3. **Hubble tension resolution:** entirely late-time, via void stress depletion:
Theoretical: $H_{\text{local}}/H_0 = 1.084 \Rightarrow 67.4 \times 1.084 = 73.1$ (SH0ES: 73.04 ± 1.04) MCMC-derived: $H_{0,\text{global}} = 68.56 + \text{void enhancement } 5.81 = H_{0,\text{local}} = 74.37$
4. **Late-time phenomenology:** $w(z)$, BAO distances, growth rate, rotation curves.
5. **Framework stability:** MTDF recovers Λ CDM at early times, deviates at late times.
6. **Cross-probe integrity (Phase 6, 10 tests):** redshift transition confirmed (3.6σ); weak

lensing and CMB lensing pipelines pass strict systematics controls; derived parameters self-consistent; S_8 tension reduced from 3.5σ to 1.9σ (combined WL); $f\sigma_8$, BAO residuals, and cluster mass predictions computed - all sensitivity-limited with current data, with quantified thresholds for future surveys.

9.5 Joint implications for late-time cosmological tensions

While MTDF is not tuned to resolve any single cosmological tension, the validated framework exhibits simultaneous shifts in multiple late-time observables. In particular, the Phase 5 best-fit favours a modest reduction in σ_8 relative to Λ CDM best-fit values (0.795 vs 0.809), consistent with the direction preferred by weak-lensing surveys (KiDS, DES).

Independently, the environment-dependent dynamics tested in Phases 3 and 4 naturally modify local expansion observables without compromising global CMB consistency. Together, these results indicate that MTDF addresses multiple late-time phenomenological tensions within a single, unified dynamical framework, without introducing additional free parameters beyond k_f .

9.6 Paper implications

Paper	Impact
MTDF_01 (Master)	Phase 1 strengthens all claims. No changes needed.
MTDF_02 (Part I, Environmental)	Signal independently confirmed (Phase 3). Redshift transition localised to $z < 0.04$ (Phase 6A, 3.6σ). Framework’s strongest empirical evidence.
MTDF_04 (Part III, Photon coupling)	$\kappa \approx f_{\text{kick}}/3$ is an observational anchor structurally related to the early field energy fraction and below Planck sensitivity. Paper frames this as speculative, which is appropriate.
MTDF_05 (Part IV, Cosmological)	Phase 5 full plik total $\Delta\chi^2 = +0.63$; Phase 2 plik-lite precursor: $+0.81$; CMB+BAO summary run: $+1.25$ (narrower data). Hubble tension mechanism fully intact. Priority 1 hard falsifier cleared.
MTDF_07 (This paper)	Phase 6 adds ten cross-probe tests (6A–6E, 6C2–6C5). Title updated to Phases 1–6. All validation claims current as of February 2026.

10 Open questions and next steps

10.1 NGC/SGC asymmetry diagnostics (Phase 3b, complete)

Phase 3b (Section 5b) implemented the full asymmetry diagnostic battery and is now complete. All five tests confirm that the NGC/SGC contrast is not statistically significant, not caused by survey geometry, and attributable to SGC void-catalogue sparsity. The NGC/SGC claim status in the taxonomy (Section 7.2) has been upgraded from $\blacktriangle$ Hypothesis to $\checkmark$ Validated.

10.2 Falsifier and discriminator tests

#	Test	Type	Status	Timescale
1	Full Planck plik + class_mtdf	Discriminator	PASSED (Phase 5, $\Delta\chi^2 = +0.63$)	Complete
2	NGC/SGC diagnostic battery	Upper limit	COMPLETE (sensitivity-limited, not geometric)	Complete
3	Weak lensing $\times$ environment (KiDS)	Upper limit	COMPLETE (Phase 6B/B2: systematics-clean null)	Complete
4	CMB lensing $\times$ voids	Upper limit	COMPLETE (Phase 6D/E: pipeline validated)	Complete
5	S ₈ cross-probe coherence	Coherence	COMPLETE (Phase 6C2: $3.5\sigma \rightarrow 1.9\sigma$)	Complete
6	$f\sigma_8 \times$ environment	Forecast	COMPLETE (Phase 6C4: 0.22σ , sensitivity-limited)	Complete
7	BAO residual structure	Forecast	COMPLETE (Phase 6C3: $\Delta W = 0.27\%$, undetectable)	Complete
8	Cluster $M_{\text{dyn}}/M_{\text{lens}}$	Resolved	RESOLVED (Phase 6C5/C5b: $\eta = 1$, Case B, $M_{\text{dyn}}/M_{\text{lens}} = 1$)	Complete
9	Weak lensing P(k) (non-linear)	Discriminator	Requires full P(k) from class_mtdf	Months
10	Merging cluster offsets	Discriminator	Requires N-body simulations	Institutional

Phase 2 scope note, resolved by Phase 5. Phase 2 used a perturbative correction layer on a Λ CDM emulator, not the full modified Boltzmann equations. Phase 5 subsequently ran `class_mtdf` through the full Planck plik likelihood with all nuisance parameters, confirming $\Delta\chi^2 = +0.63$. The perturbative approximation in Phase 2 was validated by the full calculation in Phase 5.

11 Appendix: Technical notes

11.1 A.1 k_f normalisation conventions

Three different k_f conventions were identified across the MTFD codebase during Phase 2. Understanding these is essential for interpreting historical results.

Convention	"Full MTFD" k_f	f_{kick}	Source
A (old CLASS, pre-correction)	0.10	0.003303	November 2025 cobaya runs
B (corrected CLASS)	1.0	0.003303	Theory prediction
C (Phase 2 code, initial, corrected)	0.001326	0.001326	CosmoPower correction layer (bug)

Convention C initially used $f_{\text{kick}} = \alpha\kappa = 0.001326$ instead of the correct $f_{\text{kick}} = \lambda_{\text{MTDF}}/24 = 0.003303$. This factor-of-2.49 error caused a $3\times$ overcorrection in the angular scale ratio per unit k_f , leading to a spurious H_0 - k_f degeneracy in the first MCMC run (best-fit $H_0 = 55.6$, which is unphysical). After correction, all runs show $H_0 \approx 67$ and no degeneracies.

11.2 A.2 V18 parameter reference

Parameter	Value	Uncertainty	Role
α	1.30	0.26	Stress-field coupling (void dynamics: KBC 2013 + Hoscheit 2018; counter: Kenworthy 2019)
β	22.7 Mpc	0.09×10^{23} m	Coherence length (void radius distribution: Pan+2012; Hamaus+2014)
τ	13.0 Gyr	0.2	Field age
β_{eos}	0.573	0.012	Equation-of-state parameter
E	9.1×10^{-10} Pa	0.4×10^{-10}	Field energy scale
δ_{bf}	0.150	0.030	Observational anchor: cluster baryon fraction at $M_{500} \sim 10^{14} M_{\odot}$ (Giodini+2009; Gonzalez+2013; see MTDF_01 §5.1)
z_t	0.74	0.10	Transition redshift
κ	0.00102	0.00003	Photon-stress coupling (anchor: $\kappa \approx f_{\text{kick}}/3$; below FIRAS sensitivity)
H_0	70.0 km/s/Mpc	1.5	Hubble constant (Freedman 2021 TRGB: 69.8 ± 0.6)
f_{kick}	0.003303	–	Derived: $(1 - \beta_{\text{eos}})^2 / ((1 + \alpha) \times 24)$

11.3 A.3 Historical context: November 2025 runs

Previous cobaya + class_mtdf runs (November 2025) provide useful reference points. A single-point evaluation at the Planck Λ CDM best-fit gave $\chi^2 = 614.23$ (consistent with full plik TT+TE+EE). An MCMC with Planck TT-only (16 walkers, 200 steps) found $k_f \approx 0.03$ in the old normalisation (Convention A), with χ^2 indistinguishable from Λ CDM. These early results are consistent with the current Phase 2 conclusion that Planck cannot constrain the EFE amplitude.

12 References

References

- [1] Aghanim, N. et al. (Planck Collaboration) 2020, "Planck 2018 results. V. Power spectra and likelihoods", *A&A*, 641, A5. DOI: 10.1051/0004-6361/201936386
- [2] Aghanim, N. et al. (Planck Collaboration) 2020, "Planck 2018 results. VI. Cosmological parameters", *A&A*, 641, A6. DOI: 10.1051/0004-6361/201833910
- [3] DESI Collaboration 2025, "DESI 2024 VI: Cosmological Constraints from BAO", *JCAP*, 2025, 021. DOI: 10.1088/1475-7516/2025/02/021
- [4] Brout, D. et al. 2022, "The Pantheon+ Analysis: Cosmological Constraints", *ApJ*, 938, 110. DOI: 10.3847/1538-4357/ac8e04
- [5] Melia, F. & Maier, R. S. 2013, "Cosmic chronometers in the $R_h = ct$ universe", *MNRAS*, 432, 2669. DOI: 10.1093/mnras/stt1082
- [6] Alam, S. et al. (eBOSS Collaboration) 2021, "Completed SDSS-IV eBOSS: Cosmological implications", *Phys. Rev. D*, 103, 083533. DOI: 10.1103/PhysRevD.103.083533
- [7] Spurio Mancini, A. et al. 2022, "CosmoPower: emulating cosmological power spectra for accelerated Bayesian inference", *MNRAS*, 511, 1771. DOI: 10.1093/mnras/stac064

- [8] Torrado, J. & Lewis, A. 2021, "Cobaya: Code for Bayesian Analysis of hierarchical physical models", *JCAP*, 2021, 057. DOI: 10.1088/1475-7516/2021/05/057
- [9] Foreman-Mackey, D. et al. 2013, "emcee: The MCMC Hammer", *PASP*, 125, 306. DOI: 10.1086/670067
- [10] Blas, D., Lesgourgues, J. & Tram, T. 2011, "The Cosmic Linear Anisotropy Solving System (CLASS). Part II", *JCAP*, 2011, 034. DOI: 10.1088/1475-7516/2011/07/034
- [11] Asgari, M. et al. (KiDS Collaboration) 2021, "KiDS-1000 Cosmology: Cosmic shear constraints", *A&A*, 645, A104. DOI: 10.1051/0004-6361/202039070
- [12] Abbott, T. M. C. et al. (DES Collaboration) 2022, "Dark Energy Survey Year 3 results: Cosmological constraints", *Phys. Rev. D*, 105, 023520. DOI: 10.1103/PhysRevD.105.023520
- [13] Dalal, R. et al. (HSC Collaboration) 2023, "Hyper Suprime-Cam Year 3 Results: Cosmology from Cosmic Shear", *Phys. Rev. D*, 108, 123519. DOI: 10.1103/PhysRevD.108.123519
- [14] Keenan, R. C., Barger, A. J. & Cowie, L. L. 2013, "Evidence for a ~ 300 Megaparsec Scale Under-density", *ApJ*, 775, 62. DOI: 10.1088/0004-637X/775/1/62
- [15] Pan, D. C. et al. 2012, "Cosmic Voids in SDSS DR7", *MNRAS*, 421, 926. DOI: 10.1111/j.1365-2966.2011.20197.x
- [16] Freedman, W. L. 2021, "Measurements of the Hubble Constant: Tensions in Perspective", *ApJ*, 919, 16. DOI: 10.3847/1538-4357/ac0e95
- [17] McGaugh, S. S., Lelli, F. & Schombert, J. M. 2016, "Radial Acceleration Relation in Rotationally Supported Galaxies", *PRL*, 117, 201101. DOI: 10.1103/PhysRevLett.117.201101
- [18] Lelli, F. et al. 2017, "One Law to Rule Them All: The Radial Acceleration Relation", *ApJ*, 836, 152. DOI: 10.3847/1538-4357/836/2/152
- [19] Riess, A. G. et al. 2022, "A Comprehensive Measurement of the Local Value of the Hubble Constant", *ApJL*, 934, L7. DOI: 10.3847/2041-8213/ac5c5b